

Linear Filters

Noise:

Salt & Pepper Noise: Random occurrences of black and white pixels

Impulse Noise: Random occurrences of white pixels

Gaussian Noise: Variations in intensity drawn from a Gaussian normal distribution

Filters:

Linear: Obeys Superposition, Shift-Invariant, Casual, and Stable

Non-Linear: Opposite of linear

Smoothing Filter: Values positive, $\sum F = 1$, amount of smoothing proportional to mask size, remove high-frequency components; low-pass filter

Low-Pass: Smoothing, removing noise, blurring

High-Pass: Sharpening, accentuates edges

Derivative Filter: Opposite signs to get high response in areas of high contrast, $\sum F = 0$, High absolute value when high contrast

Correlation Filter (Cross-Correlation): Pixel is linear combination of surrounding pixels, $G = H \otimes F$.

Gaussian Filter: Linear, smoothing, σ =variance, kernel=size of mask

Sharpening Filter: Accentuates differences with local average, subtraction

Convolution: Linear, Flip the filter in both dimensions (bottom to top, right to left), then apply cross-correlation, $G = H \star F$, $\frac{\partial}{\partial x}(H \star F) = (\frac{\partial}{\partial x}H) \star F$

Shift Invariant: Operator behaves the same everywhere; the value of the output depends on the pattern in the image neighborhood, not the position of the neighborhood

Superposition: The response to a sum of inputs is the sum of the responses to the individual inputs

Seperability: A 2D filter is separable if it can be written as the outer product of two 1D filters

Median Filter: Non-linear, no new pixel values, removes spikes, good for impulse + salt & pepper noise, edge preserving

Laplacian Filter: Hybrid images, Unit impulse - gaussian \approx laplacian of gaussian

Edge Detection

Edge: rapid change in image intensity, extrema of the first derivative, zero-crossings of the second derivative

Steps: 1. Smooth/Suppress noise, 2. Edge enhancement/Filter for contrast, 3. Edge localization /Local maxima/Threshold/Thinning

Canny Edge Detector:

Algorithm: 1. Filter image with derivative of Gaussian then get magnitude/direction of gradient, 2. Non-maximum suppression, 3. Linking and thresholding (hysteresis)

Property: Threshold strong edges and weak edges using

two thresholds, then keep weak edges only if connect to strong edges.

Property: Filtering a signal f with a Gaussian and then calculating its gradient is the same as filtering the signal f with the first order derivative of the Gaussian.

Seam Carving

$$\text{Energy: } \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Seam: Path of least energy, 8-connected, 1 pixel wide

Cumulative Energy Map: $M(i,j) = \text{Energy}(i,j) + \min(M(i-1,j-1), M(i-1,j), M(i-1,j+1))$

Template Matching

Template: Uses filters that match effect intended to search

Algorithm: Use normalized (Controls relative brightness) cross-correlation score to find template in image

Binary Image Analysis

Steps: 1. Threshold image into binary form, 2. Clean using morphological operations, 3. Extract separate blobs/Connected components, 4. Describe blobs with region properties

Otsu's Method: Find threshold that minimizes intra-class variance

Morphological Operations:

Dilation: Closes gaps and fills holes, $I \oplus [1]_{i \times i}$

Erosion: Erode connected components, Shrink features, Removes bridges/branches/noise, $I \ominus [1]_{i \times i}$

Opening: Erosion then dilation, removes noise/small objects, keeps original shape

Closing: Dilation then erosion, closes small holes inside or small black points on the object.

Texture

Segmentation: Analyze/represent texture, group image regions with consistent texture

Synthesis: Generate new texture patches/images

Filter Banks: Collection of (d) features, d-dimensional vector, contain combination of scales, orientation, and patterns

First-order Markov Random Field: Each pixel depends on its neighbors, $P(X|A,B,C,D)$

Single Pixel Synthesis: To synthesize x , pick one matching window at random, assign x to be the center pixel of that window, find the best matches using SSD error and randomly choose between them, preferring better matches with higher probability

Window size: Larger→More uniform, Smaller→More randomness

Image Quilting: Synthesize texture by stitching together overlapping patches from a sample image

Optical Flow

Hough Transform

Algorithm: 1. Get edges of image, 2. For each edge point vote for possible parameters (Increment accumulator array), 3. Threshold accumulator array to find matches

Pros: All points are processed independently, so can cope with occlusion/gaps, Some robustness to noise: noise points unlikely to contribute *consistently* to any single bin, Can detect multiple instances of a model in a single pass

Cons: Complexity of search time increases exponentially with the number of model parameters, Non-target shapes can produce spurious peaks in parameter space, Quantization: can be tricky to pick a good grid size

RANSAC

Description: Choose s (minimum size for fitting) samples, Fit a model to the samples, Count inliers, Repeat then take model with largest set of inliers

Adaptive RANSAC Algorithm: Let $N=\infty$, $S_{IN}=\emptyset$ and $iter=0$

While $N > iter$:

Estimate parameters a_{tst} from a random n -tuple from S

Determine inlier set S_{tst} , i.e., data points within a distance t of the model $y = f(x; a_{tst})$

If $|S_{tst}| > |S_{IN}|$:

Set $S_{IN} = S_{TST}$, $a = a_{tst}$, $w = \frac{|S_{IN}|}{|S|}$ and $N =$

$\frac{\log(1-p)}{\log(1-w^n)}$ with $p = 0.99$

end-if

iter++

end-while

Pros: Robust iterative method for estimating the parameters of a mathematical model from a set of observed data containing outliers, Separates the observed data into inliers and outliers, Can be applied in an iterative manner to obtain multiple models

Cons: Not perfect

Robust Fitting

Description: Non-linear optimization problem that must be solved iteratively, Least squares solution can be used for initialization, Adaptive choice of scale: approx. 1.5

times median residual

Active Contours-Snakes

Internal Energy:

Formula: $E_S = \alpha \left| \frac{dv}{ds} \right|^2 + \beta \left| \frac{d^2v}{ds^2} \right|^2$

Elasticity: 1st-order term, Membrane, α controls tension along spline, Stretching balloon / elastic band, Ideal curve is point

Stiffness: 2nd-order term, Thin plate, β controls rigidity of spline, Bending thin plate / wire, Ideal curve is circle

Cons: Sensitive to initial position/parameters, Small capture range, Fails to detect concavities

Level Sets

Description: Embed curve in one higher dimension; curve is given by zero level set of implicit function (i.e., intersection of function with $z=0$)

Pros: Curve may change topology and form sharp corners ('weak solutions'), Discrete grid and finite differences approximate solution, Intrinsic geometric properties are easily determined (normal vector, curvature), Formulation is same for 2D or 3D

Grouping

Bottom-Up: Pixels belong together because they look similar

Top-Down: Pixels belong together because they are from the same object

Bottom-Up Segmentation via Clustering:

Description: Separate image into coherent objects, Group together similar-looking pixels for efficiency of further processing

Algorithms: Mode finding and mean shift: k-means, EM, mean-shift, Graph-based: normalized cuts

Features: Color, texture, Quantization for texture summaries

Clustering: Unsupervised learning, Detect patterns in unlabeled data, Useful when don't know what you're looking for, Requires data, but no labels

K-means:

Algorithm: 1. Randomly initialize k cluster centers, 2. Assign each point to nearest cluster center, 3. Solve for mean of each cluster, 4. Repeat until convergence

Pros: Simple, Fast to compute, Converges to local minimum of within-cluster squared error *Cons:* Simple, Non-deterministic, Requires initial means and k , Sensitive to initial centers, Sensitive to outliers, Detects spherical clusters, Assuming means can be computed

Mean Shift:

Description: Cluster: all data points in the attraction basin of a mode, Attraction basin: the region for which

all trajectories lead to the same mode, Find features (color, gradients, texture, etc), Initialize windows at individual feature points, Perform mean shift for each window until convergence, Merge windows that end up near the same “peak” or mode

Pros: Does not assume shape on clusters, One parameter choice (window size, aka “bandwidth”), Generic technique, Find multiple modes

Cons: Selection of window size, Does not scale well with dimension of feature space

Normalized Cuts:

Pros: Generic framework, flexible to choice of function that computes weights (“affinities”) between nodes, Does not require model of the data distribution

Cons: Time complexity can be high, Preference for balanced partitions